

Computing the Formal and Institutional Boundaries of Contemporary Genre and Literary Fiction

Natasha Johnson¹ 

¹ Language Technologies Institute, Carnegie Mellon University, Pittsburgh, USA

Abstract

Though the concept of genre has been a subject of discussion for millenia, the relatively recent emergence of genre fiction has added a new layer to this ongoing conversation. While more traditional perspectives on genre have emphasized form, contemporary scholarship has invoked both formal and institutional characteristics in its taxonomy of genre, genre fiction, and literary fiction. This project explores the soundness of genre as a formal designation, using computational methods to identify the prevalence of narrative features both within specific genres and within genre versus literary fiction at large, and examining how this varies with author gender. Pulling from Andrew Piper's CONLIT dataset of Contemporary Literature, we assemble a corpus of literary and genre fiction, with the latter category containing romance, mystery, and science fiction novels. We calculate the distribution of narrative features according to author gender within each genre and within genre versus literary fiction at large. We then use logistic regression to model the effect that each feature has on literary classification and to measure how author gender moderates these effects. This project finds statistically significant formal markers of each literary category and highlights the unexpected ways in which author gender impacts literary classification.

Keywords: computational literary studies, genre

1 Introduction

Scholars have been using, debating, and reinventing the term “genre” for millennia. In fact, in 1976, Todorov admitted that discussing genre might be considered “an obviously anachronistic pastime” [8]. However, the emergence of genre fiction over the past century has added new layers to this age-old conversation.

When scholars such as Frye, Jameson, and Rosen have considered form alongside genre, they have touched on features including narrative structure, character, and the presence or lack of “elevated” or “non-elevated” language. Scholars emphasizing the institutional nature of genre have highlighted its relation to author gender and literary distribution. This project uses computational methods to further examine the formal character of genre and literary fiction and the role that author gender has on these boundaries. Through a distant reading of Andrew Piper's CONLIT corpus of contemporary fiction, we seek to gain a deeper and more empirical understanding of the formal boundaries of contemporary genre fiction, the distinctions between genre, genre fiction, and literary fiction, and how this varies by author gender.

We begin by using Welch's ANOVA to examine narrative structure, character, and elevated language between and within genre and literary fiction. Then, we use logistic regression to explore how author gender impacts the role that these features play in the classification of literature. Finally, we use unigram vectors and non-parametric distance-based statistical tests to compare the most frequently used words within each category.

Natasha Johnson. “Computing the Formal and Institutional Boundaries of Contemporary Genre and Literary Fiction.” In: *Computational Humanities Research 2025*, ed. by Taylor Arnold, Margherita Fantoli, and Ruben Ros. Vol. 3. Anthology of Computers and the Humanities. 2025, 1016–1046. <https://doi.org/10.63744/Ztw8stkvUjM1>.

© 2025 by the authors. Licensed under Creative Commons Attribution 4.0 International (CC BY 4.0).

2 Related Work

2.1 Debates on Genre

Northrop Frye is one of the earliest literary scholars whose writings on genre are still frequently brought up in current criticism. In *Anatomy of Criticism* [11], Frye presents a string of classically inspired theories about mode, symbolism, myth, and genre. His book is very abstract and distant from traditional literary scholarship, but his identification of four fundamental plot structures has been revisited by scholars such as Hayden White [39] and Kurt Vonnegut [38].

Fredric Jameson [19] builds upon Frye's work in "Magical Narratives: Romance as Genre," exploring genre both as a formal designation and as a broader social framework. Jameson proposes that genres are contracts between writers and readers that establish expectations regarding the content and form of novels and also provide a means of categorizing books according to their kind. He suggests two approaches to the practice of contemporary genre, the syntactic, which explores the form of genre and the semantic, which explores the meaning of genre within society at large.

Jeremy Rosen [32] also engages with both form and society in his discussion of genre, but he first makes the key contribution of distinguishing between "genre," "genre fiction," and "literary fiction." He defines genre as "an existing literary framework or recipe that writers may adapt according to their needs," which touches upon formal elements similar to Frye's work and Jameson's syntactic approach. Rosen suggests that authors of literary fiction make use of the same genre conventions that popular genre fiction follows, but also clearly try to create distance between their work and popular genre fiction through methods such as using elevated literary devices or critiquing popular culture within the text itself. Despite these formal practices, Rosen argues that literary fiction and genre fiction should be thought of primarily as "relational subfields of production, circulation, and reception," as opposed to formal designations. Rosen's work complements Ken Gelder's stance on genre and genre fiction. Gelder [12] suggests that genre not only serves a formal purpose for genre fiction, but that it furthermore determines a book's publishing, marketing, and consumption. Gelder emphasizes the social and industrial characteristics of genre fiction by referring to it as "popular fiction."

Mark McGurl [23] similarly challenges the validity of formal distinctions between genre and literary fiction by, in Goldstone's words, making high-literary assumptions about genre and simultaneously deflating the pretensions of literary fiction [14]. However, McGurl also suggests that digital distribution of literature undermines once strong institutional divides. Others such as Theodore Martin [22], Andrew Hoberek [17], and Marc Verboord [37] have also spoken of a narrowing gap between genre and literary fiction, but they have attributed this to a widespread elevation in the status of genre fiction and popular culture.

Formal Elements of Genre and Literary Fiction. Despite differing views on the importance of form in the classification of genre and literary fiction relative to institutional factors, scholars often identify a consistent set of formal elements to literary fiction: elegant prose, novelty, emphasis on character over plot, complex structure, serious subject matter, and slow pacing. [32; 33]. In contrast, genre fiction is often stated to be plot driven, highly standardized, fast-paced, and stylistically simple [8; 31].

Institution and Genre. Scholarship emphasizing the institutional nature of genre and literary fiction has offered varied perspectives on the specific institutions at play. Some scholars have emphasized literary fiction's resistance to mass culture [24] and distance from the marketplace [32]. Others have centered differences in the publication, distribution, and consumption of literary and genre fiction [32]. Andrew Hoberek has cited gender as the primary force at play, stating that "[g]ender continues to mark the outer limits of what counts as literary," with genre fiction often being associated with women's writing [16]. Verboord [37] has also emphasized the status-based system which underlies the literary field and the systematic marginalization of women in fields of artistic production.

2.2 Genre vs Literary Analysis

A number of empirical studies have responded to theories posed above. Algee-Hewitt et al. [1] compared canonical works (often equated with literary fiction) to archival works (other works of fiction) from 1800 - 1900 and found the canonical to have a much higher entropy but lower type token ratio than archival works. In looking at reading difficulty within books from the Chigaco Corpus (from 1880 - 2000), Bizzoni et al. [6] found that there was a small but statistically significant difference in the reading difficulty of literary vs non-literary novels.

Anderson et al. [2] analyzed 604 documents from the British National Corpus, published before 2001, and found significant differences in the linguistic and stylistic footprints male and female authors—namely that women use more pronouns, whereas men use more noun specifiers, and women are more involved, whereas men are more informational. However, it is difficult to know whether these findings apply to contemporary fiction, because the British National Corpus contains both fiction and non-fiction, and Anderson et al. observed that the aforementioned female stylistic footprints are broadly correlated with fiction, whereas the male foorprints are associated with nonfiction.

Verboord [37] looked at US, French, and German literature between 1960 and 2009 and found that women are more represented on best seller lists than among literary award winners. In line with the work of Tuchman and Fortin and Radway, he discusses how the artistic products of women are held in less artistic status than those of men [31] and how men routinely enter previously women-dominated fields that are rising in prestige and edge women out [36].

While each of these analyses offers insights into stylistic and institutional qualities of fiction, each only considers a limited range of stylistic or institutional features, without exploring ways in which these might work in combination with each other. Furthermore, none of the studies examine books written in the past two decades, after the widespread growth of the internet.

3 Data

The data we examine in this project comes from Piper’s CONLIT dataset of Contemporary Literature [30]. This dataset includes data and metadata from 2,754 Anglophone books written between 2001 and 2021 inclusive. The corpus includes 234 mystery books, 208 romance books, and 223 science fiction novels, all hand-tagged. We remove the 4 books tagged as both Science Fiction and Mystery and keep the remaining as our genre fiction books.

Category	Dataset Count
Science Fiction	222
Mystery	231
Romance	208
Literary Fiction	191
Genre Fiction	661

Table 1: Book categories explored in this project and the total number of works present within our dataset for each category

While there are no books explicitly tagged as literary fiction, the “Prizelists” category includes 258 works of fiction that have won or been shortlisted for the Faulkner Award, National Book Award, Man Booker Prize, Giller Prize, and Governor General’s Award. The former four of these awards are given to works of high literary value, which we would colloquially refer to as literary fiction; in fact, Rosen cites the selection criteria for the Man Booker Prize as standards for literary fiction [32]. The Governor General’s Award does not make this distinction and has been given to works that would generally be considered genre fiction. We remove the books nominated for this

Feature	Description
Circuitousness	Measure of narrative non-linearity
Average Speed	Measure of narrative pace
Minimum Speed	Measure of narrative distance
Event Count	Estimated number of diegetic events
Topical Heterogeneity	Measure of the semantic spread of book sections
Total Characters	Total number of named characters appearing >5 times
Protagonist Concentration	Percentage of all character mentions by main character
Tuldava Score	Reading difficulty measure
Avg. Word Length	Average length of all words per book
Avg. Sentence Length	Average length of all sentences per book
Token Count	Number of tokens (words and punctuation) in a book
Probability of First Person	1 if first-person, 0 if third person ¹

Table 2: Feature descriptions for narrative analysis

award to produce a group of 191 novels, which we consider our literary fiction. The final book counts for this project can be seen in Table 1.

For each of the books in CONLIT, Piper has shared features such as the narrative speed, topical heterogeneity, and tulvada count (a commonly used measure of reading difficulty). The features were identified using methods described by Toubia et al. [35], Ciaramita and Altün [7], and Bamman et al. [4]. The features we consider in this project are explained in Table 2. Piper also shares the top 10,000 words that appear in at least 18% of the books in the corpus, alongside their frequency counts within each book. We consider these as well.

4 Methodology

4.1 Feature comparison with Welch’s ANOVA

We begin by measuring the significance of narrative features in genre classification. We take the mean of each feature within each category (science fiction, mystery, romance, literary fiction, and genre fiction) and use Welch’s ANOVA [26] to test whether these values are significantly different between the individual genres (science fiction, mystery, romance) and between genre and literary fiction.²

For clarity, we group the features into descriptors of narrative structure, language, and character. Within narrative structure, we include narrative pace, distance, circuitousness, topical heterogeneity, and event count. Within language, we consider token count, average sentence length, average word length, Tuldava score, and the point of view (probability of 1st person). For character, we include character count and the protagonist concentration.

Then, we divide our corpus by author gender (excluding the 5 books with an unspecified author gender) and repeat the same tests for male and female authors. Book counts for male and female authors can be seen in Table 3. Since only 2 of the books written by male authors are romance books, we exclude these books as well, as they do not constitute a large enough sample size to meet the Normality assumption of Welch’s ANOVA (see Appendix A). In order to ensure that our observations about female-versus-male-authored books are not skewed by the presence of romance novels only in the former category, we run Welch’s ANOVA twice for the female authors—once with romance books included, and once with romance books excluded.

¹ For calculations involving the probability of 1st person (i.e. Welch ANOVA testing and regression modeling focusing on this feature) probability of 1st person is treated as a binary feature, and the 22 books with values in between 0 and 1 are excluded.

² For more information Welch’s ANOVA and how our data meets its assumptions, see Appendix A.

Class	Female	Male
SF	72	150
MY	117	110
ROM	205	2
Lit	80	111
Genre	394	262
Total	474	373

Table 3: Book counts within dataset by author gender

4.2 Modeling the impact of gender

To further examine the impact of author gender, we measure whether author gender affects the effect that literary features have on a book’s classification as genre or literary fiction. Using a subset of features from Table 2, we fit a logistic regression model that incorporates the main effects of each feature and author gender, as well as the interaction affects of the features and author gender. Prior to modeling, we normalize each non-binary feature using z-score standardization. The model is represented by the following formula:

$$\log \left(\frac{P(\text{Literary})}{1 - P(\text{Literary})} \right) = \beta_0 + \delta G + \sum_{i=1}^p \beta_i X_i + \sum_{i=1}^p \gamma_i (X_i \times G)$$

where

- $P(\text{Literary})$ is the probability of literary classification,
- X_i are the features,
- G is the binary gender indicator (e.g., 0 for Male, 1 for Female),
- β_0 is the intercept,
- δ is the coefficient for gender,
- β_i is the coefficient for feature i ,
- γ_i is the coefficient for the interaction between feature i and gender.

To formally test whether the interaction term significantly improves the model, we fit a second model without the interaction term, represented by the following formula:

$$\log \left(\frac{P(\text{Literary})}{1 - P(\text{Literary})} \right) = \beta_0 + \delta G + \sum_{i=1}^p \beta_i X_i$$

We then compare both models using a likelihood ratio test [27].

Feature Selection. Analysis of Variance Inflation Factors (VIFs) reveals high multicollinearity among our features (see Appendix D). To help us determine which features to exclude, we fit a logistic regression model for each feature individually, represented by the following formula:

$$\log \left(\frac{P(\text{Literary} = 1)}{1 - P(\text{Literary} = 1)} \right) = \beta_0 + \beta_i X_i + \delta G + \gamma_i (X_i \cdot G)$$

We eliminate those features with p values $> 0.05^3$. The features without significant impact are total number of characters, narrative distance, and topical heterogeneity. Then, since average sentence length and Tuldava score still have VIFs > 10 , we eliminate Tuldava score, which has the higher VIF of the two (due to its additional overlap with average word length). The remaining features

³ for results of univariate logistic regression, see Appendix E

we use for our multivariate logistic regression are protagonist concentration, average word length, average sentence length, token count, pace, circuitousness, point of view, and event count.

Probabilities. To further illustrate how author gender affects the relationship between narrative feature and literary classification, we predict the probability of literary classification given author gender and a range of values for each selected feature. We calculate these probabilities by taking the sigmoid of the log-odds (logit) for each gender along a range of values for each feature [21]. When reporting these results, we convert the scaled feature values back to their original values for easier interpretation.

4.3 Content Comparison via Unigram Analysis

We use the unigram counts as another method of measuring similarity between individual genres and between genre and literary fiction at large. This allows us to evaluate whether stereotypes about the formulaic nature of genre fiction and the novelty of literary fiction hold true.

We treat each book’s unigram vector as a feature vectors. We begin by performing PERMDISP tests with 999 permutations to understand the spread of content within each category and to explore whether the categorical variances are too large for traditional PERMANOVA tests, which require homogeneity of variance [3].

Because our PERMDISP tests reveal heteroscedasticity between the individual genres, we use W_d^* and TwT tests—non-parametric, distance-based tests of group difference and dispersion, which are robust against non-homogeneous variances [15]—to measure the distance between our unigram vectors in each category. The results of these tests can be seen in Tables 9, 10, and 11.

5 Results

5.1 Feature Welch ANOVA Testing.

Tables 4, 5, and 7 contains the means, F-statistics and effect sizes resulting from our Welch ANOVA tests.⁴ Appendix C includes histograms depicting feature usage within genre fiction and between genre and literary fiction by author gender. When looking at books written by both male and female authors, science fiction, Romance, and Mystery are significantly different ($p < 0.05$) across all the features. Likewise, Literary and Genre fiction are significantly different in protagonist concentration, sentence length, word length, Tuldava score, number of events, circuitousness, narrative distance, and topical heterogeneity. They are not significantly different in their token counts, number of characters, narrative speed, and point of view. Splitting the dataset by author gender changes the significant features.

5.2 Modeling the Effect of Author Gender

Table 8 shows the results of our logistic regression modeling. The following features have an independently significant effect ($p < 0.05$) on a book’s classification as literary or genre fiction: protagonist concentration, average word length, average sentence length, and circuitousness. Furthermore, female authorship has a significant effect on the impact of average word length in literary classification. Looking at female authorship independent of feature usage, our coefficient (-0.53) and odds ratio (0.59) suggest a potential decrease in the odds of being classified as literary compared to male authors (potentially 41% lower odds), but this potentially effect is not statistically significant at the 95% confidence interval ($p = 0.13$, 95% CI [-1.22, 0.15]). Therefore, we cannot draw definite conclusions about the effect of author gender on literary classification, independent of literary features.

⁴ For detailed results of all Welch ANOVA tests, see Appendix B

Feature	Inter-Genre					Genre vs Literary Fiction			
	SF Mean	MY Mean	ROM Mean	F	η_p^2	Genre Mean	Lit Mean	F	η_p^2
Narrative Pace	2.25	2.27	2.33	62.07*	0.15	2.28	2.29	1.32	0.00
Narrative Distance	2.48	2.50	2.56	93.61*	0.22	2.51	2.49	15.71*	0.02
Circuitousness	0.23	0.24	0.23	4.98*	0.02	0.23	0.20	38.25*	0.06
Topical Heterogeneity	3.01	3.01	3.04	6.63*	0.02	3.02	2.98	12.71*	0.02
Event Count	7286	6659	6575	3.49*	0.01	6844	5657	21.90*	0.03
Character Count	52.96	48.27	27.89	113.14*	0.15	43.43	49.25	3.51	0.01
Protagonist Conc.	0.24	0.23	0.31	34.03*	0.09	0.26	0.21	25.64*	0.03
Written in 1st Person	0.37	0.51	0.61	12.91*	0.04	0.49	0.57	3.63	0.00
Token Count	144272	124636	103573	34.41*	0.10	124604	120939	0.52	0.00
Avg Sentence Length	14.53	13.82	13.37	14.10*	0.04	13.92	16.72	113.77*	0.17
Avg Word Length	4.24	4.12	3.99	138.90*	0.32	4.12	4.17	15.99*	0.01
Tuldava Score	3.50	3.31	3.16	58.28*	0.16	3.33	3.66	123.95*	0.14

Table 4: Welch's ANOVA statistics comparing inter-genre (SF, MY, ROM) and genre vs literary fiction for combined author gender. Values are rounded to two decimal places, excluding event count and token count, which are rounded to zero decimal places. * indicates significance ($p < 0.05$). For standard deviations and degrees of freedom, see Appendix B.

Feature	Inter-Genre					Genre vs Literary Fiction			
	SF Mean	MY Mean	ROM Mean	F	η_p^2	Genre Mean	Lit Mean	F	η_p^2
Narrative Pace	2.28	2.29	2.33	16.52*	0.08	2.31	2.30	0.18	0.00
Narrative Distance	2.50	2.52	2.56	33.82*	0.14	2.54	2.50	13.28*	0.04
Circuitousness	0.22	0.23	0.23	2.91	0.01	0.23	0.19	12.90*	0.05
Topical Heterogeneity	3.02	3.03	3.04	1.31	0.01	3.03	2.99	8.06*	0.03
Event Count	7186	6742	6583	1.08	0.01	6740	5596	9.10*	0.02
Character Count	49.97	45.21	27.76	61.16*	0.24	45.65	37.00	4.62*	0.02
Protagonist Conc.	0.24	0.23	0.31	27.78*	0.12	0.27	0.22	17.83*	0.04
Written in 1st Person	0.39	0.54	0.60	4.90*	0.02	0.54	0.51	0.32	0.00
Token Count	138125	124244	103533	19.39*	0.10	116005	117351	0.04	0.00
Avg Sentence Length	14.54	13.61	13.38	7.08*	0.04	13.66	16.61	68.80*	0.18
Avg Word Length	4.18	4.08	3.99	49.00*	0.22	4.05	4.15	37.72*	0.06
Tuldava Score	3.44	3.25	3.16	20.40*	0.11	3.24	3.64	100.36*	0.19

Table 5: Welch's ANOVA statistics comparing inter-genre (SF, MY, ROM) and genre vs literary fiction for female authors. Values are rounded to two decimal places, excluding event count and token count, which are rounded to zero decimal places. * indicates significance ($p < 0.05$). For standard deviations and degrees of freedom see Appendix B.

Feature	Inter-Genre				Genre vs Literary Fiction			
	SF Mean	MY Mean	<i>F</i>	η_p^2	Genre Mean	Lit Mean	<i>F</i>	η_p^2
Token Count	138125	124244	4.34*	0.03	129532	117351	2.92	0.01
Character Count	49.97	45.21	1.75	0.01	47.03	45.65	0.11	0.00
Protagonist Conc.	0.24	0.23	0.87	0.00	0.23	0.22	1.90	0.01
Avg Sentence Length	14.54	13.61	7.69*	0.04	13.96	16.61	49.77*	0.19
Avg Word Length	4.18	4.08	19.47*	0.10	4.12	4.15	2.95	0.01
Tuldava Score	3.44	3.25	15.72*	0.08	3.32	3.64	53.16*	0.17
Event Count	7186	6742	1.30	0.01	6911	5596	11.20*	0.05
Narrative Pace	2.28	2.29	0.02	0.00	2.29	2.30	2.70	0.01
Circuitousness	0.22	0.23	5.25*	0.03	0.23	0.19	11.22*	0.06
Narrative Distance	2.50	2.52	4.68*	0.02	2.51	2.50	1.90	0.01
Topical Heterogeneity	3.02	3.03	0.47	0.00	3.02	2.99	5.54*	0.03
Written in 1st Person	0.39	0.54	4.21*	0.02	0.48	0.51	0.16	0.00

Table 6: Welch’s ANOVA statistics comparing inter-genre (SF & MY) and genre vs literary fiction for female authors. Values are rounded to two decimal places, except for token and event count, which are rounded to zero decimal places. * indicates significance ($p < 0.05$). For standard deviations and degrees of freedom, see Appendix B.

Feature	Inter-Genre				Genre vs Literary Fiction			
	SF Mean	MY Mean	<i>F</i>	η_p^2	Genre Mean	Lit Mean	<i>F</i>	η_p^2
Token Count	147223	125185	10.92*	0.04	137899	123525	3.71	0.01
Character Count	54.39	51.36	0.63	0.00	53.11	51.85	0.07	0.00
Protagonist Conc.	0.24	0.23	0.42	0.00	0.24	0.21	4.64*	0.01
Avg Sentence Length	14.53	14.08	2.62	0.01	14.34	16.80	42.09*	0.14
Avg Word Length	4.27	4.15	45.41*	0.14	4.22	4.18	5.99*	0.01
Tuldava Score	3.53	3.37	15.54*	0.05	3.47	3.67	22.97*	0.06
Event Count	7335	6595	4.57*	0.01	7022	5702	13.87*	0.04
Narrative Pace	2.24	2.24	0.11	0.00	2.24	2.28	19.23*	0.05
Circuitousness	0.23	0.25	7.99*	0.03	0.24	0.20	31.53*	0.09
Narrative Distance	2.47	2.49	6.02*	0.02	2.48	2.48	0.22	0.00
Topical Heterogeneity	3.00	2.99	0.52	0.00	3.00	2.98	1.85	0.01
Written in 1st Person	0.36	0.49	4.46*	0.02	0.42	0.61	13.20*	0.03

Table 7: Welch’s ANOVA statistics comparing inter-genre (SF & MY) and genre vs literary fiction for male authors. Values are rounded to two decimal places, except for token and event count, which are rounded to zero decimal places. * indicates significance ($p < 0.05$). For standard deviations and degrees of freedom, see Appendix B.

The results of our Likelihood Ratio test ($\chi^2(7) = 20.12$, $p = 0.0053$) confirm that including the 7 author-feature interaction terms in our model significantly improves the fit of the model when compared to the model without any interaction terms.

Parameter	Coefficient	P-Value	CI Lower	CI Upper	Odds Ratio
Intercept	-1.23	0.00	-1.74	-0.73	0.29
Female Author	-0.53	0.13	-1.22	0.15	0.59
Protagonist Concentration	-0.36	0.02	-0.66	-0.06	0.70
+ Female Author	-0.16	0.52	-0.66	0.33	0.85
Average Word Length	-0.69	0.00	-1.08	-0.30	0.50
+ Female Author	0.76	0.01	0.16	1.35	2.13
Average Sentence Length	1.11	0.00	0.76	1.47	3.05
+ Female Author	-0.03	0.90	-0.54	0.48	0.97
Token Count	0.19	0.47	-0.32	0.69	1.20
+ Female Author	0.65	0.14	-0.20	1.49	1.91
Pace	0.18	0.34	-0.20	0.56	1.20
+ Female Author	0.10	0.74	-0.46	0.65	1.10
Circuitousness	-0.58	0.00	-0.96	-0.20	0.56
+ Female Author	-0.03	0.93	-0.59	0.54	0.97
Written in 1st Person	0.05	0.88	-0.62	0.72	1.05
+ Female Author	-0.24	0.64	-1.25	0.76	0.78
Event Count	-0.40	0.15	-0.94	0.14	0.67
+ Female Author	-0.48	0.26	-1.33	0.36	0.62

Table 8: Logistic regression results with interaction terms, rounded to two decimal places. For standard errors and z-statistics, see Appendix E

Figure 1 illustrates the probability of having one’s book classified as literary fiction based on the author gender and feature usage, for the features which have a significant effect ($p < 0.05$) on literary classification. Average sentence length is positively correlated with literary classification, whereas circuitousness and protagonist concentration are inversely correlated with literary classification. For male authors, average word length is inversely correlated with literary classification, whereas for female authors, the correlation between average word length and literary classification is positive.

5.3 Unigram Comparison

Table 10 shows the results of the PERMDISP tests assessing the dispersion within individual genres and within genre and literary fiction. The individual genres are significantly different in their dispersions ($p=0.0040$), and thus it would not be possible to compare them using a method that assumes homogeneity of variance. The broader category of literary and genre fiction are not significantly different in their variances.

Table 11 includes the results of our W_d^* and T_w^2 tests, showing the distances between the means of the unigram vectors from within each category. We see that literary and genre fiction are much more distinctive in their content and word usage than the individual genres are. Within genre pairings, romance and science fiction are the most distinct, and science fiction and mystery are the least distinct. The p-values indicate that all these values are significant ($p < 0.05$), meaning that none of these distinctions could have occurred by chance.

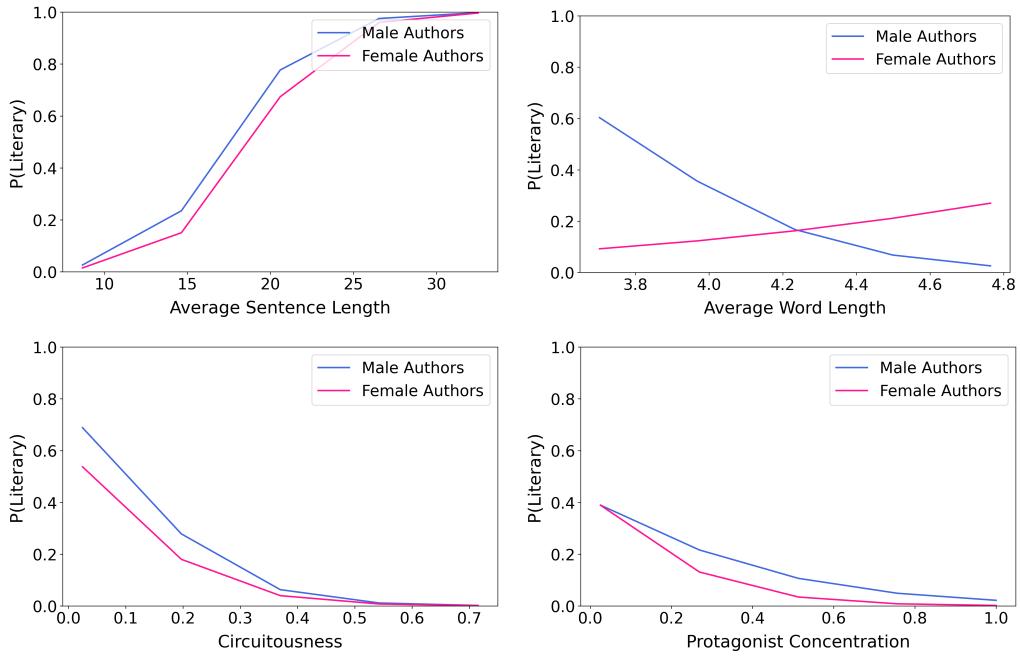


Figure 1: Plots depicting the probability of literary classification relative to feature usage, for features with a significant impact ($p < 0.05$) on model results. Note that only the interaction between author gender and average word length is significant ($p < 0.05$) at a 95% confidence interval.

Category	Variance
SF	0.0032
MY	0.0060
ROM	0.0045
Genre	0.0055
Literary	0.0043
Genre and Literary	0.0061

Table 9: Variance within and across categories, rounded to four decimal places.

6 Discussion

Our analyses reveals key formal characteristics of individual genres and of genre and literary fiction at large, while also highlighting how author gender affects writing practices and literary classification.

6.1 Characteristics of Individual Genres

Each of the genres examined demonstrates strong conventions with regards to narrative structure, characters, and style.

Contemporary mystery novels feature more circuitous plots than science fiction or romance novels. Such non-linearity has long been considered a characteristic element of mystery novels, due to the fact that they intertwine two stories—that of the crime and that of its investigation [18]. Our mystery novels also feature the lowest protagonist concentration of the genre fiction categories, meaning that the protagonist makes up the lowest percentage of total character mentions. This cannot be attributed to a higher number of total characters or to a lower number of first-person narrators

Test Comparison	F-Statistic	P-Value	Mean Distance to Centroid
Genre Fiction (SF, MY, ROM)	5.0330	0.0040	MY: 0.0437 ROM: 0.0403 SF: 0.0341 Genre: 0.0447
Literary vs Genre	1.0058	0.2813	Lit: 0.0420

Table 10: Results of PERMDISP tests with 999 repetitions. Values have been rounded to four decimal places.

Test	P-Value	Wasserstein Statistic
W_d^* Multi-group test	0.00	75.87
Lit vs Genre T_w^2	0.00	163.86
MY vs ROM T_w^2	0.00	50.57
MY vs SF T_w^2	0.00	19.34
ROM vs SF T_w^2	0.00	147.36

Table 11: Results of W_d^* and T_w^2 tests with 999 repetitions. Values have been rounded to two decimal places.

within mystery novels than within other genres (as neither of those are true). It may be reflective of the unimportance of character within the mystery genre: Nearly a century ago, when speaking of the leisurely enjoyment of mystery novels by academics, Nicolson said, "Character...troubles us little, though characters we have in abundance" [28]. More recently, Goldman has spoken of the widely critiqued one-dimensionality of characters within mystery novels [13] (though he highlights specific examples of characters who he believes go against this stereotype). Whether due to reader preferences, genre expectations, or the former affecting the latter, our results suggest that these claims about character still hold true for contemporary works in the genre.

Romance novels cover more distance than mystery or science fiction, meaning that their ending states are the most removed from their starting states. This intuitively makes sense given that most romance novels end with a happy couple (which is not how they often begin). Romance novels also have the fastest pace of all the genres and cover the most topics in each book section. Interestingly, fewer actual events occur in romance novels than in science fiction or mystery novels. One potential explanation for this is how protagonist-centered romance novels are: they feature the protagonist most out of any genre, they feature fewer side characters than other genres, and they are more likely to be written in first-person. It seems likely that romance novels prioritize revealing protagonist interiority over describing external events. Furthermore, romance novels are by far the easiest genres to read, as indicated by their Tuldava scores, average word lengths, and average sentence lengths. Some have asserted that romance is not a respected because of its association with women [10], and while many scholars have made strong arguments for this [31; 36; 37], its syntactic simplicity might also be a contributing factor.

Overall, science fiction features the most events but covers the least narrative distance. It is significantly slower paced than romance. Science fiction narratives feature more characters than other genres and are least likely to be written in first-person. Though the protagonist concentration within science fiction is slightly higher than in mystery novels, it is much lower than in romance novels. Similarity to with mystery novels, this could be attributed to a lack of emphasis on character, which has long-been discussed as a common characteristic of science fiction [5]. Regardless of author gender, science fiction texts are consistently the most difficult genre to read as indicated by their Tuldava score, average word length, and average sentence length, and total token count.

These observations highlight the formulaic nature of genre fiction, in line with existing schol-

arly literature [8; 31]. Furthermore, our analysis of unigram vectors shows a distinctiveness of content between genres, in line with many theories of genre which view genres as "regions of conceptual space" [9].⁵

Gendered Differences within Genres.

One of the most noteworthy observations about gendered practices within genre fiction begins with our corpus: Out of the 207 romance novels in the CONLIT dataset, only 2 were written by men. Though female authors have long been the primary producers of romance fiction [37], this staggering of a difference far exceeds author gender ratios suggested by other researchers [34].⁶ The difference could be attributed to the fact that the romance novels in the CONLIT dataset were drawn from a list of bestselling romance novels on Amazon, in which female authors might be more highly represented than in the romance genre more broadly.

While romance novels are distinct from science fiction and mystery novels both in terms of their form and in terms of their author gender, science fiction and mystery texts are quite similar across female and male authors. For the 6 features which are significantly different between genres for authors of both genders and for the 2 features which are significantly different between genres for one gender, the genre with the higher mean value is consistent across gender.

6.2 Genre versus Literary Fiction

The broader categories of genre and literary fiction exhibit noteworthy distinctions with regards to form, though not necessarily in the areas proposed by existing scholarship [8; 31; 32; 33]. Compared to genre fiction, literary fiction does contain longer and more difficult to read sentences, which supports the notion that literary fiction features more elevated language.

In contrast with the idea that literary fiction typically features slow, complex narratives, our analysis finds that literary fiction is not significantly slower-paced than genre fiction overall and is consistently more linear in its narratives. In fact, increased circuitousness (narrative non-linearity) is a significant predictor of genre fiction as opposed to literary fiction, as illustrated in Figure 1. Thus, the supposed elevated and complex character of literary fiction do not extend beyond language into more complex narrative structures.

Prestige by Author Gender.

When looking at science fiction, mystery, and literary fiction, the difference in feature distribution between genre and literary fiction is significant in 8 categories for male authors but only 5 categories for female authors. Though there are fewer categories of difference between genre and literary fiction for female authors, the differences in some of these categories is larger for female authors than for male. For example, our Welch's ANOVA results show that Tuldava score and average sentence length have higher effect sizes for female authors than for male authors. It seems possible that the threshold for literary status for female authors is dependent on fewer integrated textual characteristics and is more heavily determined by superficial textual elevatedness. This would be supported by the results of our logistic regression modeling, which reveal that longer words are positively correlated with literary classification for female authors but are negatively correlated with literary classification for male authors (see Figure 1). Together, these observations suggest that the formal boundaries between genre and literary fiction are poorly defined, particularly for female authors.

⁵ We cite Evnine's work because we use his exact wording to refer to a collection of theories of genre, not because he argues in favor of such theories.

⁶ It is important to note that while WordsRated identifies as a literary research group, their findings are not externally validated or peer-reviewed

6.3 Limitations

Narrative Features. Our approaches were limited by the fact that the corpus does not include the full texts for each book. While we were able to perform basic similarity comparison using the unigram vectors, word-count vectors are not able to capture the same level of nuance and meaning that contextualized embeddings would be able to achieve. Furthermore, due to our not having access to the full bodies of the text, our analyses largely depend on the pre-extracted literary features included in the CONLIT dataset. Though the methods for extracting these features were validated by the researchers who developed them, they were not all evaluated on contemporary fiction, which might impact their effectiveness.

Author Gender. The books in the CONLIT dataset were manually labeled for author gender based on the perceived gender of the author. There is a chance that some of the authors published under pen names associated with the opposite gender [25], which would have resulted in their gender being inaccurately represented within the dataset and within this project.

7 Conclusion

Our project highlights both the formal and institutional character of genre. Our combined-gender science fiction, mystery, and romance groups all have distinctive formal fingerprints, and many of the features that contribute to these fingerprints are features that have long-been associated with these genres by academics and non-academics alike.

Our analyses similarly reveal differences in the forms of literary and genre fiction. Across both male and female authors, longer sentences and decreased circuitousness and protagonist concentrations are predictive of literary classification. However, the use of longer words is positively correlated with literary status for female authors but is negatively correlated with literary status for male authors. Results such as these illustrate the complex and unexpected ways in which author gender interacts with form when literary status is being determined.

Acknowledgements

I would like to thank Jeffrey Williams for his literature recommendations and feedback on the early stages of this project; the CHR reviewers for their comments and suggestions; Jake Hofgard for the helpful discussions on statistical modeling; and Andrew Piper for assembling and sharing the CONLIT dataset.

References

- [1] Algee-Hewitt, Mark, Allison, Sarah, Gemma, Marissa, Heuser, Ryan, Moretti, Franco, Walser, Hannah, and Lab, Literary. *Canon/Archive. Large-scale Dynamics in the Literary Field*. Jan. 2016. URL: <https://litlab.stanford.edu/LiteraryLabPamphlet11.pdf>.
- [2] Anderson, Marti J. “A New Method for Non-parametric Multivariate Analysis of Variance”. In: *Austral Ecology* 26, no. 1 (Feb. 2001), pp. 32–46. DOI: 10.1111/j.1442-9993.2001.01070.pp.x. URL: <https://doi.org/10.1111/j.1442-9993.2001.01070.pp.x>.
- [3] Anderson, Marti J. “Distance-Based Tests for Homogeneity of Multivariate Dispersions”. In: *Biometrics* 62, no. 1 (2006), pp. 245–253. DOI: 10.1111/j.1541-0420.2005.00440.x. URL: <https://doi.org/10.1111/j.1541-0420.2005.00440.x>.

- [4] Bamman, David, Underwood, Ted, and Smith, Noah A. “A Bayesian Mixed Effects Model of Literary Character”. In: *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics Volume 1: Long Papers* (Jan. 2014). DOI: 10.3115/v1/p14-1035. URL: <https://doi.org/10.3115/v1/p14-1035>.
- [5] Bereit, Virginia F. “The Genre of Science Fiction”. In: *Elementary English* 46, no. 7 (1969), pp. 895–900. ISSN: 00135968. URL: <http://www.jstor.org/stable/41386588> (visited on 10/28/2025).
- [6] Bizzoni, Yuri et al. “Good Reads and Easy Novels: Readability and Literary Quality in a Corpus of US-published Fiction”. Tech. rep. 2023, pp. 42–51. URL: <https://aclanthology.org/2023.nodalida-1.5.pdf>.
- [7] Ciaramita, Massimiliano and Altün, Yasemin. “Broad-coverage sense disambiguation and information extraction with a supersense sequence tagger”. In: *Proceedings of the 2006 Conference on Empirical Methods in Natural Language Processing* (Jan. 2006). DOI: 10.3115/1610075.1610158. URL: <https://doi.org/10.3115/1610075.1610158>.
- [8] Duff, David. *Modern Genre Theory*. Routledge, Jan. 2000.
- [9] Evnine, Simon J. “”But is It Science Fiction?”: Science Fiction and a Theory of Genre”. In: *Midwest Studies in Philosophy* 39, no. 1 (2015), pp. 1–28. DOI: 10.1111/misp.12037.
- [10] Fendley, Gina. “Is Fanfiction Better than Published Romance Novels?” Tech. rep. 2019. URL: <https://marymount.edu/academics/wp-content/uploads/sites/3/2023/05/Magnificat-2019.pdf#page=127>.
- [11] Frye, Northrop. *Anatomy of Criticism*. Dec. 1957. DOI: 10.1515/9781400866908. URL: <https://doi.org/10.1515/9781400866908>.
- [12] Gelder, Ken. *Popular Fiction: The Logistics and Practices of a Literary Field*. Routledge, Dec. 2004.
- [13] Goldman, Alan H. “The Appeal of the Mystery”. In: *Journal of Aesthetics and Art Criticism* 69, no. 3 (Aug. 2011), pp. 261–272. DOI: 10.1111/j.1540-6245.2011.01470.x. URL: <https://doi.org/10.1111/j.1540-6245.2011.01470.x>.
- [14] Goldstone, Andrew B. “Genre Fiction Without Shame”. In: *American literary history* 35, no. 4 (Nov. 2023), pp. 1745–1758. DOI: 10.1093/alh/ajad152. URL: <https://doi.org/10.1093/alh/ajad152>.
- [15] Hamidi, Bashir, Wallace, Kristin, Vasu, Chenthamarakshan, and Alekseyenko, Alexander V. “W*d-test: Robust Distance-based Multivariate Analysis of Variance”. In: *Microbiome* 7, no. 1 (Apr. 2019). DOI: 10.1186/s40168-019-0659-9. URL: <https://doi.org/10.1186/s40168-019-0659-9>.
- [16] Hoberek, Andrew. *American Literature in Transition, 2000–2010*. Cambridge University Press eBooks, Dec. 2017. DOI: 10.1017/9781316569290. URL: <https://doi.org/10.1017/9781316569290>.
- [17] Hoberek, Andrew. “Introduction: After Postmodernism”. In: *Twentieth Century Literature* 53, no. 3 (Sept. 2007), pp. 233–247. DOI: 10.1215/0041462x-2007-4007. URL: <https://doi.org/10.1215/0041462x-2007-4007>.
- [18] Hühn, Peter. “THE DETECTIVE AS READER: NARRATIVITY AND READING CONCEPTS IN DETECTIVE FICTION”. In: *Modern Fiction Studies* 33, no. 3 (1987), pp. 451–466. ISSN: 00267724, 1080658X. URL: <http://www.jstor.org/stable/26282385> (visited on 10/28/2025).

- [19] Jameson, Fredric. "Magical Narratives: Romance as Genre". In: *New Literary History* 7, no. 1 (Jan. 1975), p. 135. DOI: 10.2307/468283. URL: <https://doi.org/10.2307/468283>.
- [20] Kwak, Sang Gyu and Kim, Jong Hae. "Central Limit Theorem: The Cornerstone of Modern Statistics". In: *Korean journal of anesthesiology* 70, no. 2 (Jan. 2017), p. 144. DOI: 10.4097/kjae.2017.70.2.144. URL: <https://doi.org/10.4097/kjae.2017.70.2.144>.
- [21] Li, Hongyu and Barry, Jay. "Interpreting Interactions in Logistic Regression". Tech. rep. Cornell Statistical Consulting Unit, Apr. 2022.
- [22] Martin, Theodore. *Contemporary Drift*. Literature Now, Jan. 2017.
- [23] McGurl, Mark. *Everything and Less*. Verso Books, Oct. 2021.
- [24] McGurl, Mark. *The Novel Art*. Princeton University Press, Nov. 2001.
- [25] McLean, Madeleine. "Men With Female Pen Names". 2023. URL: <https://sites.duke.edu/unsuitable/men-with-female-pen-names/>.
- [26] Moder, Karl. "Alternatives to F-Test in One Way ANOVA in Case of Heterogeneity of Variances (a Simulation Study)". In: *Psychological Test and Assessment Modeling* 52 (Apr. 2010), pp. 343–353. URL: https://www.psychologie-aktuell.com/fileadmin/download/ptam/4-2010_20101218/01_Moder.pdf.
- [27] Newsom. "Interactions With Logistic Regression". In: *Psy 525/625 Categorical Data Analysis* (2021). URL: https://web.pdx.edu/~newsomj/cdaclass/ho_interactions.pdf.
- [28] Nicolson, Marjorie. *The Professor and the Detective*. GROSSET & DUNLAP, 1946, p. 116.
- [29] Piovesana, Adina and Senior, Graeme. "How Small Is Big: Sample Size and Skewness". In: *Assessment* 25, no. 6 (Sept. 2016), pp. 793–800. DOI: 10.1177/1073191116669784. URL: <https://doi.org/10.1177/1073191116669784>.
- [30] Piper, Andrew. "The CONLIT Dataset of Contemporary Literature". In: *Journal of open humanities data* 8 (Jan. 2022). DOI: 10.5334/johd.88. URL: <https://doi.org/10.5334/johd.88>.
- [31] Radway, Janice A. *Reading the romance*. Univ of North Carolina Press, Jan. 1991.
- [32] Rosen, Jeremy and Rosen, Jeremy. "Literary Fiction and the Genres of Genre Fiction". July 2019. URL: <https://post45.org/2018/08/literary-fiction-and-the-genres-of-genre-fiction>.
- [33] Saricks, Joyce G. *The Readers' Advisory Guide to Genre Fiction*. American Library Association, Apr. 2009.
- [34] Talbot, Dean. "Author Demographics Statistics". 2023. URL: <https://wordsrated.com/author-demographics-statistics/>.
- [35] Toubia, Olivier, Berger, Jonah, and Eliashberg, Jehoshua. "How Quantifying the Shape of Stories Predicts Their Success". In: *Proceedings of the National Academy of Sciences of the United States of America* 118, no. 26 (June 2021). DOI: 10.1073/pnas.2011695118. URL: <https://doi.org/10.1073/pnas.2011695118>.
- [36] Tuchman, Gaye and Fortin, Nina E. "Fame and Misfortune: Edging women out of the great literary tradition". In: *American Journal of Sociology* 90, no. 1 (July 1984), pp. 72–96. DOI: 10.1086/228048. URL: <https://doi.org/10.1086/228048>.

- [37] Verboord, Marc. “Female bestsellers: A cross-national study of gender inequality and the popular–highbrow culture divide in fiction book production, 1960–2009”. In: *European Journal of Communication* 27, no. 4 (Dec. 2012), pp. 395–409. DOI: 10 . 1177 / 0267323112459433. URL: <https://doi.org/10.1177/0267323112459433>.
- [38] Vonnegut, Kurt. “Shape of Stories”. Feb. 2004. URL: https://www.youtube.com/watch?v=G0Gru_4z1Vc.
- [39] White, Hayden. “The Historical Text as Literary Artifact”. In: *CLIO* 3, no. 3 (1971). URL: <https://www.proquest.com/docview/1300303656?pq-origsite=gscholar&fromopenview=true&sourcetype=Scholarly%20Journals&imgSeq=1>.

A Assumptions of Welch’s ANOVA

For each Welch ANOVA, our book categories are the independent variables, and our feature is the dependent variable. All of the tests are one-way. A p -score < 0.05 indicates a significant difference.

Welch ANOVA tests rely on assumptions of Independence, Additivity, and Normality [26]. Our data meets the Independence assumption because no book’s narrative speed, Tuldava score, etc. has the ability to impact other books’ measurements in those areas. Furthermore, each book is only sampled one time. The Additivity assumption is met because we have removed books tagged in multiple categories. Because the smallest sample size we consider is 191, we can assume that the sampling distribution of any of these samples is normally distributed according to the Central Limit Theorem [20], thus meeting the Normality assumption [29].

B Additional ANOVA Output

Feature	SF Mean	MY Mean	ROM Mean	SF STD	MY STD	ROM STD	df1	df2	F	p	η_p^2
Narrative Pace	2.25	2.27	2.33	0.08	0.08	0.08	2	438.01	62.07	0.00	0.15
Narrative Distance	2.48	2.50	2.56	0.06	0.06	0.06	2	436.45	93.61	0.00	0.22
Circuitousness	0.23	0.24	0.23	0.05	0.05	0.05	2	437.18	4.98	0.01	0.02
Topical Heterogeneity	3.01	3.01	3.04	0.11	0.08	0.09	2	427.92	6.63	0.00	0.02
Event Count	7286	6659	6575	3361	1843	2849	2	399.33	3.49	0.03	0.01
Character Count	52.96	48.27	27.89	36.85	18.60	13.09	2	404.03	113.14	0.00	0.15
Protagonist Conc.	0.24	0.23	0.31	0.12	0.09	0.11	2	427.13	34.03	0.00	0.09
Written in 1st Person	0.37	0.51	0.61	0.48	0.50	0.49	2	437.12	12.91	0.00	0.04
Token Count	144272	124636	103573	61930	36214	42327	2	417.69	34.41	0.00	0.10
Avg Sentence Length	14.53	13.82	13.37	2.47	2.03	2.08	2	433.98	14.10	0.00	0.04
Avg Word Length	4.24	4.12	3.99	0.17	0.13	0.15	2	429.10	138.90	0.00	0.32
Tuldava Score	3.50	3.31	3.16	0.37	0.28	0.29	2	430.46	58.28	0.00	0.16

Table 12: Welch's ANOVA statistics for inter-genre comparison with combined author gender. Values are rounded to two decimal places, except for df1, event count, and token count, which are rounded to zero decimal places. $p < 0.05$ indicates significance.

Feature	SF Mean	MY Mean	ROM Mean	SF STD	MY STD	ROM STD	df1	df2	F	p	η_p^2
Token Count	138125	124244	103533	50615	32159	42599	2	174.17	19.39	0.00	0.10
Character Count	49.97	45.21	27.76	27.60	16.62	13.13	2	149.01	61.16	0.00	0.24
Protagonist Conc.	0.24	0.23	0.31	0.09	0.09	0.11	2	190.24	27.78	0.00	0.12
Avg Sentence Length	14.54	13.61	13.38	2.30	2.09	2.08	2	174.75	7.08	0.00	0.04
Avg Word Length	4.18	4.08	3.99	0.15	0.13	0.15	2	178.83	49.00	0.00	0.22
Tuldava Score	3.44	3.25	3.16	0.34	0.28	0.29	2	172.19	20.40	0.00	0.11
Event Count	7186	6742	6583	3039	1659	2863	2	176.10	1.08	0.34	0.01
Narrative Pace	2.28	2.29	2.33	0.07	0.08	0.08	2	181.15	16.52	0.00	0.08
Circuitousness	0.22	0.23	0.23	0.04	0.05	0.05	2	183.81	2.91	0.06	0.01
Narrative Distance	2.50	2.52	2.56	0.05	0.06	0.06	2	190.72	33.82	0.00	0.14
Topical Heterogeneity	3.02	3.03	3.04	0.08	0.07	0.09	2	189.49	1.31	0.27	0.01
Written in 1st Person	0.39	0.54	0.60	0.49	0.50	0.49	2	180.02	4.90	0.01	0.02

Table 13: Welch’s ANOVA statistics for inter-genre comparison with for female authors. Values are rounded to two decimal places, except for df1, event count, and token count, which are rounded to zero decimal places. $p < 0.05$ indicates significance.

Feature	SF Mean	MY Mean	ROM Mean	SF STD	MY STD	ROM STD	df1	df2	F	p	η_p^2
Token Count	147223	125185	–	66641	40462	–	1	250.07	10.92	0.00	0.04
Character Count	54.39	51.36	–	40.55	20.06	–	1	230.00	0.63	0.43	0.00
Protagonist Conc.	0.24	0.23	–	0.13	0.09	–	1	257.47	0.42	0.52	0.00
Avg Sentence Length	14.53	14.08	–	2.56	1.98	–	1	257.18	2.62	0.11	0.01
Avg Word Length	4.27	4.15	–	0.17	0.12	–	1	257.89	45.41	0.00	0.14
Tuldava Score	3.53	3.37	–	0.39	0.27	–	1	257.15	15.54	0.00	0.05
Event Count	7335	6595	–	3515	2030	–	1	245.86	4.57	0.03	0.01
Narrative Pace	2.24	2.24	–	0.07	0.08	–	1	229.78	0.11	0.74	0.00
Circuitousness	0.23	0.25	–	0.05	0.04	–	1	249.87	7.99	0.01	0.03
Narrative Distance	2.47	2.49	–	0.06	0.06	–	1	233.16	6.02	0.01	0.02
Topical Heterogeneity	3.00	2.99	–	0.12	0.08	–	1	257.97	0.52	0.47	0.00
Written in 1st Person	0.36	0.49	–	0.48	0.50	–	1	229.27	4.46	0.04	0.02

Table 14: Welch’s ANOVA statistics for inter-genre comparison for male authors. Values are rounded to two decimal places, except for df1, event count, and token count, which are rounded to zero decimal places. $p < 0.05$ indicates significance.

Feature	Genre Mean	Lit Mean	Genre STD	Lit STD	df1	df2	F	p	η_p^2
Token Count	124604	120939	50722	64443	1	261.7	0.52	0.47	0
Character Count	43.43	49.25	27.27	40.35	1	242.26	3.51	0.06	0.01
Protagonist Conc.	0.26	0.21	0.11	0.11	1	303.73	25.64	0.00	0.03
Avg Sentence Length	13.92	16.72	2.25	3.42	1	239.31	113.77	0.00	0.17
Avg Word Length	4.12	4.17	0.18	0.14	1	393.99	15.99	0.00	0.01
Tuldava Score	3.33	3.66	0.35	0.37	1	291.39	123.95	0.00	0.14
Event Count	6844	5657	2760	3175	1	278.22	21.90	0.00	0.03
Narrative Pace	2.28	2.29	0.09	0.08	1	310.31	1.32	0.25	0
Circuitousness	0.23	0.20	0.05	0.07	1	242.14	38.25	0.00	0.06
Narrative Distance	2.51	2.49	0.07	0.08	1	274.38	15.71	0.00	0.02
Topical Heterogeneity	3.02	2.98	0.09	0.12	1	253.63	12.71	0.00	0.02
Written in 1st Person	0.49	0.57	0.50	0.49	1	313.99	3.63	0.06	0

Table 15: Welch's ANOVA statistics for literary vs genre comparison with combined author gender. Values are rounded to two decimal places, except for df1, event count, and token count, which are rounded to zero decimal places. $p < 0.05$ indicates significance.

Feature	Genre Mean	Lit Mean	Genre STD	Lit STD	df1	df2	F	p	η_p^2
Token Count	137899.33	123525.25	57990.48	68844.40	1	179.73	3.71	0.06	0.01
Character Count	53.11	51.85	33.43	43.88	1	166.92	0.07	0.79	0.00
Protagonist Conc.	0.24	0.21	0.12	0.12	1	213.29	4.64	0.03	0.01
Avg Sentence Length	14.34	16.80	2.34	3.69	1	148.98	42.09	0.00	0.14
Avg Word Length	4.22	4.18	0.16	0.15	1	226.86	5.99	0.02	0.01
Tuldava Score	3.47	3.67	0.35	0.40	1	185.84	22.97	0.00	0.06
Event Count	7022.07	5701.61	2996.06	3181.16	1	197.08	13.87	0.00	0.04
Narrative Pace	2.24	2.28	0.08	0.08	1	196.49	19.23	0.00	0.05
Circuitousness	0.24	0.20	0.05	0.06	1	171.27	31.53	0.00	0.09
Narrative Distance	2.48	2.48	0.06	0.07	1	184.41	0.22	0.64	0.00
Topical Heterogeneity	3.00	2.98	0.10	0.12	1	183.10	1.85	0.18	0.01
Written in 1st Person	0.42	0.61	0.49	0.48	1	213.53	13.20	0.00	0.03

Table 16: Welch's ANOVA statistics for genre vs literary fiction comparison for male authors. Values are rounded to two decimal places, except for df1, event count, and token count, which are rounded to zero decimal places. $p < 0.05$ indicates significance.

Feature	Genre Mean	Lit Mean	Genre STD	Lit STD	df1	df2	F	p	η_p^2
Token Count	117351.35	116004.65	58017.75	43600.76	1	97.90	0.04	0.84	0.00
Character Count	45.65	37.00	34.83	20.10	1	89.97	4.62	0.03	0.02
Protagonist Conc.	0.22	0.27	0.11	0.10	1	109.14	17.83	0.00	0.04
Avg Sentence Length	16.61	13.66	3.03	2.16	1	95.95	68.80	0.00	0.18
Avg Word Length	4.15	4.05	0.13	0.16	1	135.95	37.72	0.00	0.06
Tuldava Score	3.64	3.24	0.33	0.31	1	108.95	100.36	0.00	0.19
Event Count	5595.65	6740.27	3184.51	2604.72	1	101.54	9.10	0.00	0.02
Narrative Pace	2.30	2.31	0.09	0.08	1	106.98	0.18	0.67	0.00
Circuitousness	0.19	0.23	0.09	0.05	1	89.47	12.90	0.00	0.05
Narrative Distance	2.50	2.54	0.09	0.06	1	94.58	13.28	0.00	0.04
Topical Heterogeneity	2.99	3.03	0.13	0.08	1	92.20	8.06	0.01	0.03
Written in 1st Person	0.51	0.54	0.50	0.50	1	113.40	0.32	0.57	0.00

Table 17: Welch's ANOVA statistics for genre vs literary fiction comparison for female authors. Values are rounded to two decimal places, except for df1, event count, and token count, which are rounded to zero decimal places. $p < 0.05$ indicates significance.

C Feature Histograms

D Variance Inflation Factors

Feature	VIF
Const	2.65
Token Count	23.96
Total Characters	2.88
Protagonist Concentration	1.41
Average Sentence Length	24.65
Average Word Length	11.64
Tuldava Score	43.27
Event Count	4.61
Speed Avg	4527199.22
Circuitousness	1867745.32
Speed Min	3408000.66
Volume	37.04
Written in 1st Person	1.57

Table 18: Variance Inflation Factors (VIFs) for features, rounded to two decimal places




Figure 2: Character histograms for genre fiction by author gender

E Univariate Logistic Regression Output




Figure 3: Narrative structure histograms for genre fiction by author gender




Figure 4: Language histograms for genre fiction by author gender




Figure 5: Character histograms for literary vs genre fiction by author gender

Feature	Interaction Coeff	Interaction P-Val	LR Interaction P-Val
Circuitousness	0.80	0.82	0.81
Pace	7.07	0.00	0.00
Distance	8.98	0.00	0.00
Event Count	0.00	0.87	0.87
Topical Heterogeneity	3.52	0.04	0.04
Number of Characters	0.01	0.01	0.01
Protagonist Concentration	-3.85	0.03	0.03
Tuldava Scores	2.00	0.00	0.00
Average Word Length	5.90	0.00	0.00
Average Sentence Length	0.13	0.07	0.07
Token Count	0.00	0.15	0.15
Books in 1st Person	-1.15	0.00	0.00

Table 19: Output of univariate regression, rounded to two decimal places.




Figure 6: Narrative structure histograms for literary vs genre fiction by author gender




Figure 7: Language histograms for literary vs genre fiction by author gender




Figure 8: Character histograms




Figure 9: Narrative structure histograms




Figure 10: Language histograms